

Loop interchange

Troels Henriksen

Based on material by Cosmin Oancea

Determining parallel loops from direction vectors

Parallel loop theorem

A loop in a loop nest is parallel if all its directions are either = or there exists an outer loop whose direction is <.

Determining parallel loops from direction vectors

Parallel loop theorem

A loop in a loop nest is parallel if all its directions are either = or there exists an outer loop whose direction is <.

```
for (int i = 0; i < N; i++)  
  for (int j = 0; j < N; j++)  
    S1: A[j][i] = A[j][i]...
```

[=,=]

Determining parallel loops from direction vectors

Parallel loop theorem

A loop in a loop nest is parallel if all its directions are either = or there exists an outer loop whose direction is <.

```
#pragma omp parallel for  
for (int i = 0; i < N; i++)  
    for (int j = 0; j < N; j++)  
        S1: A[j][i] = A[j][i]...
```

[=,=]

Determining parallel loops from direction vectors

Parallel loop theorem

A loop in a loop nest is parallel if all its directions are either = or there exists an outer loop whose direction is <.

```
#pragma omp parallel for
for (int i = 0; i < N; i++)                                     [=,=]
    for (int j = 0; j < N; j++)
        S1 : A[j][i] = A[j][i]...
```



```
for (int i = 1; i < N; i++)
    for (int j = 0; j < N; j++)                                  [<,>]
        S1 : A[i][j] = A[i-1][j+1]...
```

Determining parallel loops from direction vectors

Parallel loop theorem

A loop in a loop nest is parallel if all its directions are either = or there exists an outer loop whose direction is <.

```
#pragma omp parallel for
for (int i = 0; i < N; i++)
    for (int j = 0; j < N; j++)
        S1: A[j][i] = A[j][i]...
```

[=,=]

```
for (int i = 1; i < N; i++)
    #pragma omp parallel for
    for (int j = 0; j < N; j++)
        S1: A[i][j] = A[i-1][j+1]...
```

[<,>]

- Direction vectors are not just useful for figuring out where we can add OpenMP pragmas to get free performance.
- They can also be used to tell us which *loop transformations* are valid.

Interchange

```
for (int j = 0; i < N; i++)  
  for (int i = 0; i < N; i++) ⇒ for (int i = 0; i < N; i++)  
    A[i][j] = i*j;                for (int j = 0; i < N; i++)  
                                   A[i][j] = i*j;
```

- Improves spatial locality in this case.
- Can we argue with direction vectors that this is a valid transformation?

Interchange

```
for (int j = 0; i < N; i++)  
  for (int i = 0; i < N; i++) ⇒ for (int i = 0; i < N; i++)  
    A[i][j] = i*j;                for (int j = 0; i < N; i++)  
                                    A[i][j] = i*j;
```

- Improves spatial locality in this case.
- Can we argue with direction vectors that this is a valid transformation?

Validity of loop interchange

A permutation of the loops in a loop nest is legal if permuting the direction matrix in the same way does not result in a $>$ direction as the leftmost non- $=$ direction in a row.

The loop above has direction vectors $[=,=]$ for access to A , so interchange is allowed.

A useful corollary

Interchanging a parallel loop inwards

It is always safe to interchange a *parallel loop* inwards one step at a time (i.e., if the parallel loop is the k^{th} loop in the nest then one can always interchange it with loop $k + 1$, then with loop $k + 2$, etc.).

- This means that if somehow know (by eyeballing, or an existing OpenMP pragma we trust) that a loop is parallel, we can interchange it inwards.
- We don't need to calculate the direction vectors.
- **Even if we don't actually want to execute a loop in parallel, the independence of its iterations can still be exploited for locality optimisations via interchange.**

What about loop B?

```
for (int i = 1; i < N; i++)  
  for (int j = 1; j < N; j++) {  
    S1: A[j][i] = A[j-1][i-1]...  
    S2: B[j][i] = B[j-1][i]...  
  }
```

$$M = \begin{cases} [<, <] \\ [=, <] \end{cases}$$

What about loop B?

```
for (int i = 1; i < N; i++)  
  for (int j = 1; j < N; j++) {  
    S1: A[j][i] = A[j-1][i-1]...  
    S2: B[j][i] = B[j-1][i]...  
  }
```

$$M = \begin{cases} [<, <] \\ [=, <] \end{cases}$$

```
for (int j = 1; j < N; j++)  
  for (int i = 1; i < N; i++) {  
    S1: A[j][i] = A[j-1][i-1]...  
    S2: B[j][i] = B[j-1][i]...  
  }
```

$$M = \begin{cases} [<, <] \\ [<, =] \end{cases}$$

After interchange the (now) innermost loop is parallel!

What about loop B?

```
for (int i = 1; i < N; i++)  
    for (int j = 1; j < N; j++) {  
        S1: A[j][i] = A[j-1][i-1]...  
        S2: B[j][i] = B[j-1][i]...  
    }
```

$$M = \begin{cases} [<, <] \\ [=, <] \end{cases}$$

```
for (int j = 1; j < N; j++)  
    #pragma omp parallel for  
    for (int i = 1; i < N; i++) {  
        S1: A[j][i] = A[j-1][i-1]...  
        S2: B[j][i] = B[j-1][i]...  
    }
```

$$M = \begin{cases} [<, <] \\ [<, =] \end{cases}$$

After interchange the (now) innermost loop is parallel!

What about loop C?

```
for (int i = 1; i < N; i++)  
  for (int j = 0; j < N; j++)  
     $S_1: A[i][j] = A[i-1][j+1] \dots$ 
```

$M = [<, >]$

What about loop C?

```
for (int i = 1; i < N; i++)  
  for (int j = 0; j < N; j++)  
    S1: A[i][j] = A[i-1][j+1]...
```

$M = [<, >]$

```
for (int j = 0; j < N; j++)  
  for (int i = 1; i < N; i++)  
    S1: A[i][j] = A[i-1][j+1]...
```

$M = [>, <]$

Interchange is illegal, because direction vectors are not allowed to start with >.

Interchange to increase parallelism

```
#pragma omp parallel for  
for (int i = 0; i < n; i++)  
    A[i] = f(A[i]);
```

- A common pattern in data analysis is that we have a parallel loop.

Interchange to increase parallelism

```
for (int j = 0; j < m; j++)  
    #pragma omp parallel for  
    for (int i = 0; i < n; i++)  
        A[i] = f(A[i]);
```

- A common pattern in data analysis is that we have a parallel loop.
- This loop is then contained in a sequential *time series loop*.

Interchange to increase parallelism

```
#pragma omp parallel for
for (int l = 0; l < k; l++)
    for (int j = 0; j < m; j++)
        #pragma omp parallel for
        for (int i = 0; i < n; i++)
            A[j][i] = f(A[j][i]);
```

- A common pattern in data analysis is that we have a parallel loop.
- This loop is then contained in a sequential *time series loop*.
- And this whole thing is then applied to multiple sets simultaneously.

Interchange to increase parallelism

```
#pragma omp parallel for
for (int l = 0; l < k; l++)
    for (int j = 0; j < m; j++)
        #pragma omp parallel for
        for (int i = 0; i < n; i++)
            A[j][i] = f(A[j][i]);
```

- A common pattern in data analysis is that we have a parallel loop.
- This loop is then contained in a sequential *time series loop*.
- And this whole thing is then applied to multiple sets simultaneously.
- **But wait**—usually OpenMP will only parallelise the outer k -iteration loop.

Interchange will fix this

```
#pragma omp parallel for
for (int l = 0; l < k; l++)
    for (int j = 0; j < m; j++)
        #pragma omp parallel for
        for (int i = 0; i < n; i++)
            A[j][i] = f(A[j][i]);
```

- Since the outermost loop is parallel, we can interchange it *inwards*.

Interchange will fix this

```
for (int j = 0; j < m; j++)  
    #pragma omp parallel for  
    for (int l = 0; l < k; l++)  
        #pragma omp parallel for  
        for (int i = 0; i < n; i++)  
            A[j][i] = f(A[j][i]);
```

- Since the outermost loop is parallel, we can interchange it *inwards*.
- And then we can collapse the two OpenMP directives.

Interchange will fix this

```
for (int j = 0; j < m; j++)  
    #pragma omp parallel for collapse(2)  
    for (int l = 0; l < k; l++)  
        for (int i = 0; i < n; i++)  
            A[j][i] = f(A[j][i]);
```

- Since the outermost loop is parallel, we can interchange it *inwards*.
- And then we can collapse the two OpenMP directives.
- And now we have a parallel loop with $k*n$ iterations, which is run m times.

Is that better? Depends.

Summary

- Direction vectors...
 - ▶ tell us whether a loop is parallel.
 - ▶ tell us whether loop interchange is valid.
- Loop interchange...
 - ▶ can change the locality of the loop nest (for better or worse).
 - ▶ can reveal more possibilities for parallelisation.